

# Summer 2023: ML 5710

(Assignment 2)

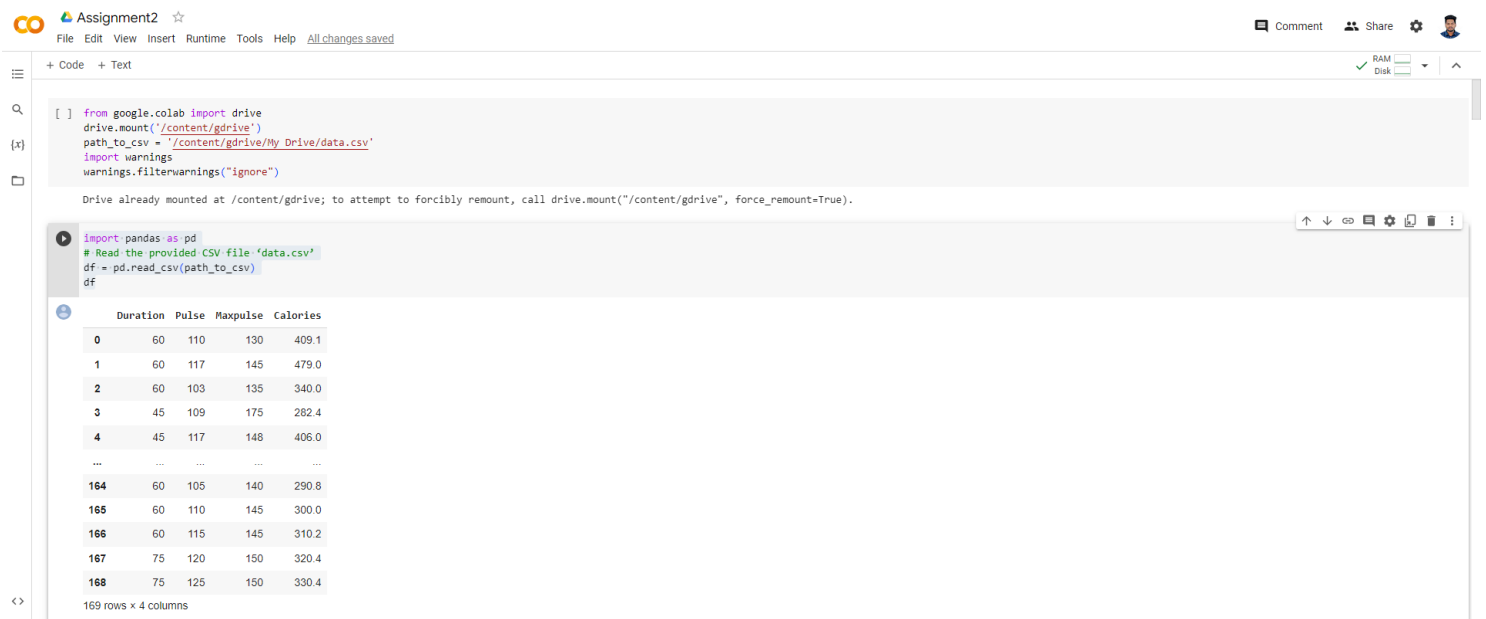
**GitHub Link:** <https://github.com/arunkumar2601/Assignment2.git>  
**Yeddula Arun Kumar** **700739780**

## Programming elements: Classification

### A.Pandas

1. Read the provided CSV file 'data.csv'.

<https://drive.google.com/drive/folders/1h8C3mLsso-R-sIOLsvoYwPLzy2fJ4IOF?usp=sharing>



The screenshot shows a Google Colab notebook titled "Assignment2". The code cell contains the following Python code:

```
[ ] from google.colab import drive
drive.mount('/content/gdrive')
path_to_csv = '/content/gdrive/My Drive/data.csv'
import warnings
warnings.filterwarnings("ignore")

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True).

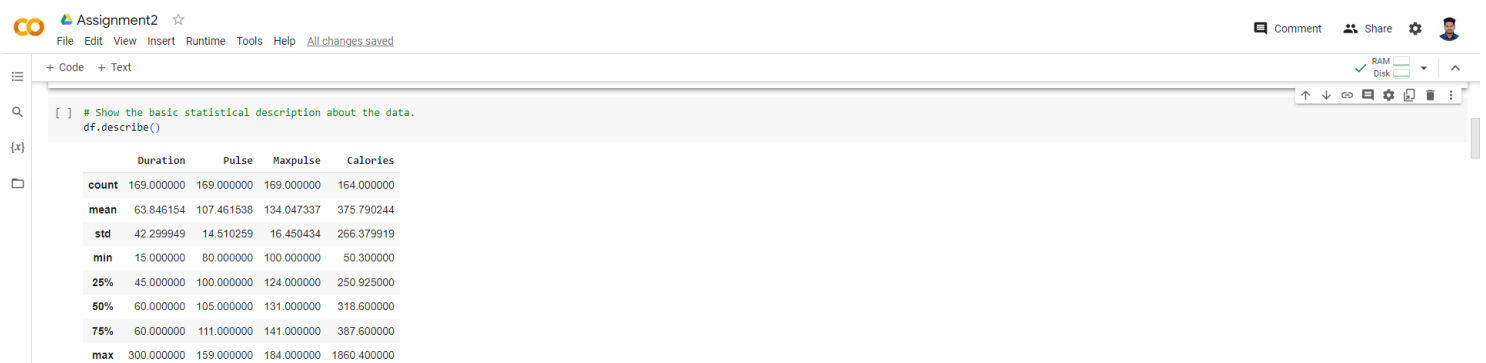
import pandas as pd
# Read the provided CSV file 'data.csv'
df = pd.read_csv(path_to_csv)
df
```

The output of the code is a pandas DataFrame with 169 rows and 4 columns: Duration, Pulse, Maxpulse, and Calories. The first few rows are displayed:

	Duration	Pulse	Maxpulse	Calories
0	60	110	130	409.1
1	60	117	145	479.0
2	60	103	135	340.0
3	45	109	175	282.4
4	45	117	148	406.0
...	...	...	...	...
164	60	105	140	290.8
165	60	110	145	300.0
166	60	115	145	310.2
167	75	120	150	320.4
168	75	125	150	330.4

The output also indicates that there are 169 rows and 4 columns in the DataFrame.

2. Show the basic statistical description about the data.



The screenshot shows the same Google Colab notebook. The code cell now contains the following Python code:

```
[ ] # Show the basic statistical description about the data.
df.describe()
```

The output of the code is a pandas DataFrame with 8 rows and 4 columns: Duration, Pulse, Maxpulse, and Calories. The rows represent statistical measures: count, mean, std, min, 25%, 50%, 75%, and max. The values are as follows:

	Duration	Pulse	Maxpulse	Calories
count	169.000000	169.000000	169.000000	164.000000
mean	63.846154	107.461538	134.047337	375.790244
std	42.299949	14.510259	16.450434	266.379919
min	15.000000	80.000000	100.000000	50.300000
25%	45.000000	100.000000	124.000000	250.925000
50%	60.000000	105.000000	131.000000	318.600000
75%	60.000000	111.000000	141.000000	387.600000
max	300.000000	159.000000	184.000000	1860.400000

3. Check if the data has null values. a. Replace the null values with the mean.

```
[ ] # Check if the data has null values.
print('Are there any null values: ',df.isnull().values.any())
# Replace the null values with the mean
df.fillna(df.mean(),inplace=True)
print('Are there any null values after using fillna: ',df.isnull().values.any())
```

Are there any null values: True  
Are there any null values after using fillna: False

4. Select at least two columns and aggregate the data using: min, max, count, mean.

```
# Select at least two columns and aggregate the data using: min, max, count, mean.
aggre = df.groupby('Duration').agg(['Calories':['mean','min','max','count']])
aggre
```

	Calories			
	mean	min	max	count
Duration				
15	87.350000	50.5	124.2	2
20	151.600000	50.3	229.4	9
25	244.200000	244.2	244.2	1
30	192.125000	86.2	319.2	16
45	279.096585	100.7	406.0	35
60	341.046465	215.2	486.0	79
75	325.400000	320.4	330.4	2
80	643.100000	643.1	643.1	1
90	541.800000	466.4	700.0	8
120	666.833333	500.0	1000.1	3
150	939.400000	816.0	1115.0	4
160	943.700000	853.0	1034.4	2
180	733.600000	600.1	800.4	3
210	1618.200000	1376.0	1860.4	2
270	1729.000000	1729.0	1729.0	1
300	1500.200000	1500.2	1500.2	1

5. Filter the data frame to select the rows with calories values between 500 and 1000

```
# Filter the dataframe to select the rows with calories values between 500 and 1000
df[(df['Calories']>=500) & (df['Calories']<=1000)]
```

	Duration	Pulse	Maxpulse	Calories
51	80	123	146	643.1
62	160	109	135	853.0
65	180	90	130	800.4
66	150	105	135	873.4
67	150	107	130	816.0
72	90	100	127	700.0
73	150	97	127	953.2
75	90	98	125	563.2
78	120	100	130	500.4
83	120	100	130	500.0
90	180	101	127	600.1
99	90	93	124	604.1
101	90	90	110	500.0
102	90	90	100	500.0
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3

## 6. Filter the dataframe to select the rows with calories values > 500 and pulse <100

# Filter the dataframe to select the rows with calories values > 500 and pulse < 100  
df[(df['Calories']>500) & (df['Pulse']<100)]

	Duration	Pulse	Maxpulse	Calories
65	180	90	130	800.4
70	150	97	129	1115.0
73	150	97	127	953.2
75	90	98	125	563.2
99	90	93	124	604.1
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3

## 7. Create a new “df\_modified” dataframe that contains all the columns from df except for “Maxpulse”.

Assignment2 ☆  
File Edit View Insert Runtime Tools Help All changes saved

Comment Share

+ Code + Text

RAM  
Disk

# Create a new “df\_modified” dataframe that contains all the columns from df except for “Maxpulse”  
df\_modified = df[['Duration', 'Pulse', 'Calories']]  
df\_modified

	Duration	Pulse	Calories
0	60	110	409.1
1	60	117	479.0
2	60	103	340.0
3	45	109	282.4
4	45	117	406.0
...	...	...	...
164	60	105	290.8
165	60	110	300.0
166	60	115	310.2
167	75	120	320.4
168	75	125	330.4

169 rows × 3 columns


## 8. Delete the “Maxpulse” column from the main df dataframe

# Delete the “Maxpulse” column from the main df dataframe  
df = df.drop('Maxpulse', axis=1)  
df

	Duration	Pulse	Calories
0	60	110	409.1
1	60	117	479.0
2	60	103	340.0
3	45	109	282.4
4	45	117	406.0
...	...	...	...
164	60	105	290.8
165	60	110	300.0
166	60	115	310.2
167	75	120	320.4
168	75	125	330.4

169 rows × 3 columns

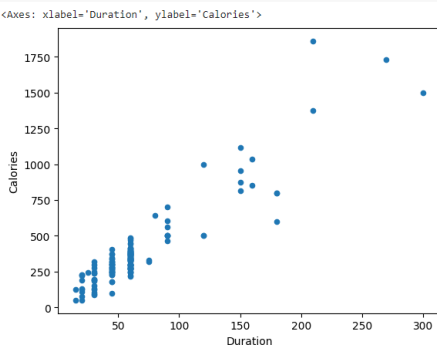
9. Convert the datatype of Calories column to int datatype.
10. Using pandas create a scatter plot for the two columns (Duration and Calories).
  - a. Example:



```
# Convert the datatype of Calories column to int datatype
df['Calories'] = df['Calories'].astype('int64')
df.dtypes
```

```
Duration    int64
Pulse       int64
Calories    int64
dtype: object
```

```
# Using pandas create a scatter plot for the two columns (Duration and Calories)
df.plot.scatter(x='Duration', y='Calories')
```



Titanic Dataset


```
[ ] import matplotlib.pyplot as plt
import seaborn as sns
train_df = pd.read_csv('/content/gdrive/My Drive/Dataset/train.csv')
test_df = pd.read_csv('/content/gdrive/My Drive/Dataset/test.csv')
titanic = pd.concat([train_df, test_df])
# Find the correlation between 'survived' (target column) and 'sex' column for the Titanic use case inclass
titanic_copy = titanic.copy()
titanic_copy['Survived'] = titanic_copy['Survived'].astype('category').cat.codes
titanic_copy['Sex'] = titanic_copy['Sex'].astype('category').cat.codes
titanic_copy['Embarked'] = titanic_copy['Embarked'].astype('category').cat.codes
corr = titanic_copy['Survived'].corr(titanic_copy['Sex'])
corr
```

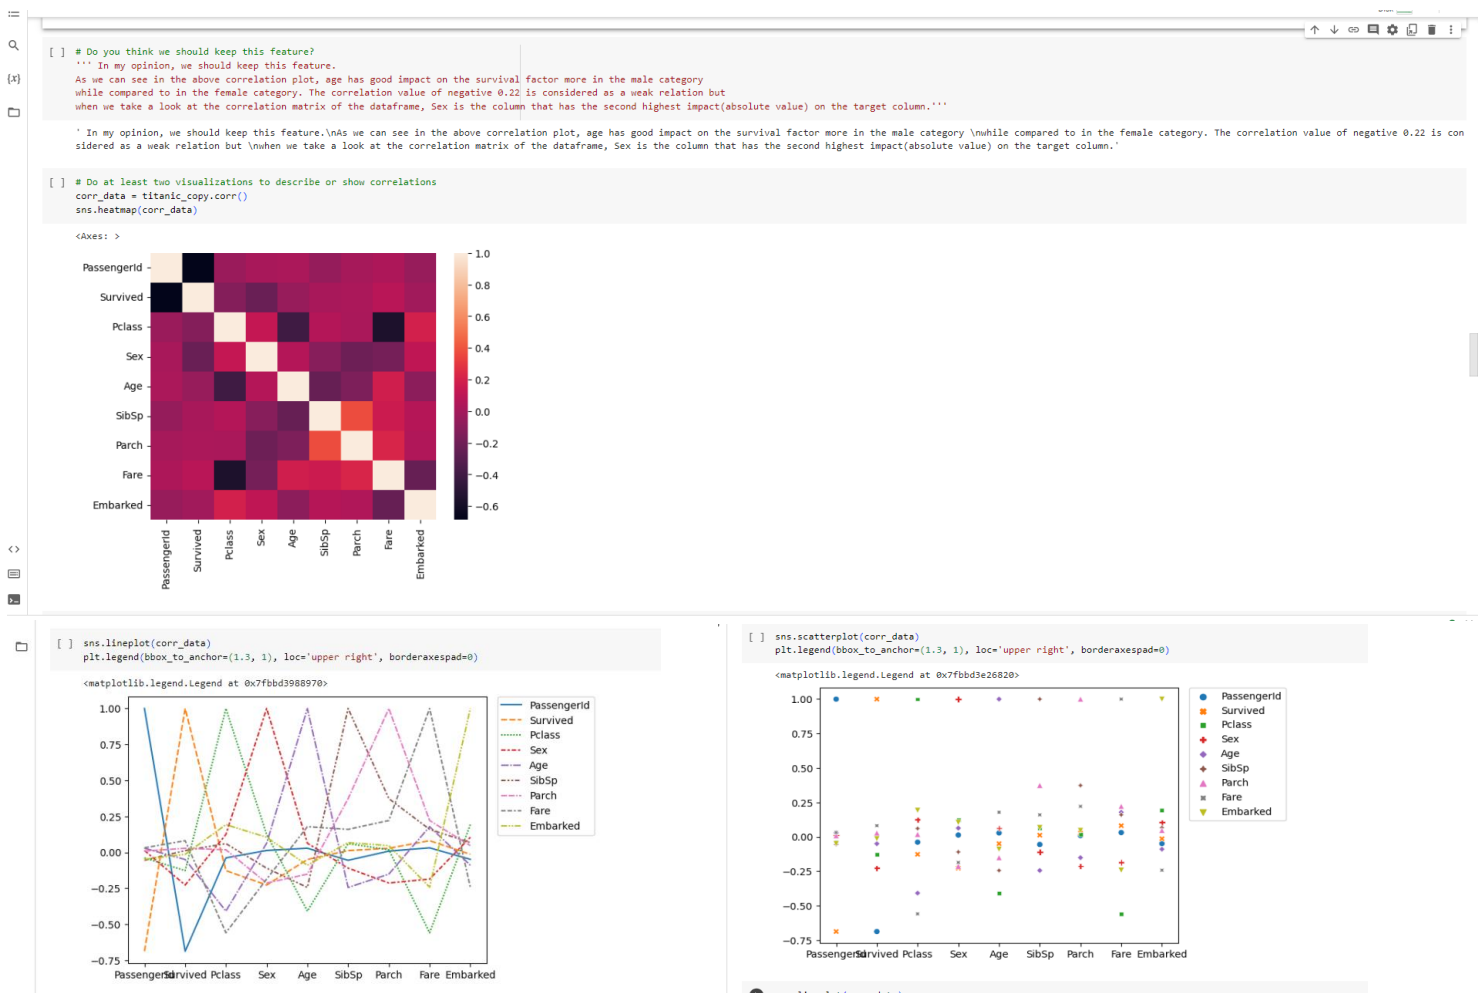
```
-0.2269178484034977
```

```
[ ] titanic_copy.corr()['Survived']
```

```
PassengerId    -0.697207
Survived        1.000000
Pclass         -0.126769
Sex            -0.226918
Age            -0.049620
SibSp          0.012470
Parch          0.028196
Fare           0.001545
Embarked       -0.013303
Name: Survived, dtype: float64
```

```
survived = 'survived'
not_survived = 'not survived'
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(10, 4))
women = titanic[titanic['Sex']=='female']
men = titanic[titanic['Sex']=='male']
ax = sns.distplot(women[women['Survived']==1].Age.dropna(), bins=18, label = survived, ax = axes[0], kde =False)
ax = sns.distplot(women[women['Survived']==0].Age.dropna(), bins=40, label = not_survived, ax = axes[0], kde =False)
ax.legend()
ax.set_title('Female')
ax = sns.distplot(men[men['Survived']==1].Age.dropna(), bins=18, label = survived, ax = axes[1], kde = False)
ax = sns.distplot(men[men['Survived']==0].Age.dropna(), bins=40, label = not_survived, ax = axes[1], kde = False)
ax.legend()
_ = ax.set_title('Male')
```





## B. Scikit-learn

1. Implement Naïve Bayes method using scikit-learn library.
  - a. Use the glass dataset available in [Link](#) also provided in your assignment.
  - b. Use **train\_test\_split** to create training and testing part.
2. Evaluate the model on testing part using score and

```
classification_report(y_true, y_pred)
```

1. Implement linear SVM method using scikit library
  - a. Use the glass dataset available in [Link](#) also provided in your assignment.
  - b. Use **train\_test\_split** to create training and testing part.
2. Evaluate the model on testing part using score and

```
classification_report(y_true, y_pred)
```

Do at least two visualizations to describe or show correlations in the Glass Dataset. Which algorithm you got better accuracy? Can you justify why?

```
+ Code + Text

# Implement Naive Bayes method using scikit-learn library and report the accuracy
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import confusion_matrix
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score
titanic_copy = titanic_copy.dropna()
df = titanic_copy.drop(['PassengerId', 'Name', 'Ticket', 'Fare', 'Cabin'], axis=1)
X = df.drop('Survived', axis=1)
y = df['Survived']
# Use train_test_split to create training and testing part.
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0)
classifier = GaussianNB()
classifier.fit(X_train, y_train)

y_pred = classifier.predict(X_test)
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
print('accuracy is', accuracy_score(y_pred, y_test))

precision    recall  f1-score   support

-1      0.75      0.16      0.26      19
 0      0.44      0.44      0.44       9
 1      0.55      0.85      0.67      27

accuracy          0.55      55
macro avg      0.58      0.48      0.46      55
weighted avg   0.60      0.55      0.49      55

[[ 3  2 14]
 [ 0  4  5]
 [ 1  3 23]]
accuracy is 0.5454545454545454
```

```
+ Code + Text

Glass dataset

# Implement Naive Bayes method using scikit-learn library.
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
glass = pd.read_csv('/content/gdrive/My Drive/Dataset/glass.csv')
X = glass.iloc[:, :-1].values
y = glass.iloc[:, -1].values
# Use train_test_split to create training and testing part.
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 0)
classifier = GaussianNB()
classifier.fit(X_train, y_train)

y_pred = classifier.predict(X_test)

# Evaluate the model on testing part using score and classification_report
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
print('accuracy is', accuracy_score(y_pred, y_test), '\n\n')
print(classification_report(y_test, y_pred))

accuracy is 0.37209302325581395

precision    recall  f1-score   support

1      0.19      0.44      0.27       9
2      0.33      0.16      0.21      19
3      0.33      0.20      0.25       5
5      0.00      0.00      0.00       2
6      0.67      1.00      0.80       2
7      1.00      1.00      1.00       6

accuracy          0.37      43
macro avg      0.42      0.47      0.42      43
weighted avg   0.40      0.37      0.36      43
```

```
# Implement linear SVM method using scikit library
from sklearn.svm import SVC

classifier_1 = SVC()
classifier_1.fit(X_train, y_train)

y_pred_1 = classifier_1.predict(X_test)
# Evaluate the model on testing part using score and classification_report
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
print('accuracy is', accuracy_score(y_pred_1, y_test), '\n\n')
print(classification_report(y_test, y_pred_1))

accuracy is 0.20930232558139536

precision    recall  f1-score   support

1      0.21      1.00      0.35       9
2      0.00      0.00      0.00      19
3      0.00      0.00      0.00       5
5      0.00      0.00      0.00       2
6      0.00      0.00      0.00       2
7      0.00      0.00      0.00       6

accuracy          0.21      43
macro avg      0.03      0.17      0.06      43
weighted avg   0.04      0.21      0.07      43

[ ] # Do at least two visualizations to describe or show correlations in the Glass Dataset.
correlation_data = glass.corr()
sns.heatmap(correlation_data)

<Axes: >
```

